

CAPSTONE PROJECT

Data Exploration and Modelling Documentation

Shreya Jain

Contents

| | |
|--|----|
| Introduction..... | 3 |
| Problem Statement | 3 |
| Motivation..... | 4 |
| Objectives | 4 |
| Project Overview | 5 |
| The features..... | 5 |
| Workspace..... | 7 |
| Importing file using pandas | 7 |
| Dataset information..... | 7 |
| Exploratory Data Analysis | 8 |
| Checking missing values..... | 8 |
| Checking unique values | 9 |
| Correlation Analysis | 9 |
| Data Distribution..... | 11 |
| Data Visualization..... | 14 |
| Data Cleaning and Preprocessing | 23 |
| Splitting the dataset..... | 24 |
| Model building..... | 25 |
| Random Forest (variable/feature importance)..... | 25 |
| Random Forest Classifier | 27 |
| Decision Tree..... | 28 |
| Logistic Regression | 32 |
| XG Boost..... | 36 |
| Model Comparison..... | 37 |
| Interpretation and Conclusion | 38 |
| Recommendations..... | 40 |

Introduction

The World Health Organization estimates that heart disease causes 12 million deaths worldwide each year. Numerous studies have been carried out to identify the most important risk factors for heart disease and to precisely estimate the overall risk. Heart disease is also referred to as a silent killer because it causes a person to pass away without any evident signs. To avoid complications in high-risk patients and make decisions about lifestyle changes, early detection of heart disease is crucial. By examining patient data that uses machine learning algorithms to categorise whether a patient has heart disease or not, this study hopes to predict future cases of heart disease.

According to estimates, 17.9 million people die from heart disease each year, accounting for 32% of all fatalities worldwide. With the help of data, we can investigate and identify the factors that most influence a person's probability of developing heart disease.

Problem Statement

The main difficulty with heart disease is detecting it. There are tools that can forecast heart disease, but they are either expensive or ineffective at calculating the likelihood of heart disease in humans. The mortality rate and total consequences can be reduced by early identification of heart disorders. Since it takes more intelligence, time, and knowledge, it is not always possible to accurately monitor patients every day, and a doctor cannot consult with a patient for a whole 24 hours. Since there is a lot of data available nowadays, we can use a variety of machine learning methods to search for hidden patterns. In medical data, the hidden patterns might be used for health diagnosis.

Motivation

For many different types of data science applications, machine learning approaches have been around, compared, and used for analysis. This research-based effort was primarily driven by the desire to investigate the methodologies for feature selection, data preparation, and processing used in machine learning training models. The difficulty we have today with first-hand models and libraries is data, which, in addition to being plentiful and our cooked models, has a bigger variation than the accuracy we observe during training, testing, and actual validation. Therefore, the goal of this study is to investigate the mechanisms underlying the models that will be used to train the gathered data.

Objectives

The following are the key goals for creating this project:

1. To implement Logistic Regression in a machine learning model to predict the likelihood of heart disease in the future.
2. To identify significant risk factors for heart disease based on medical data.
3. To improve their chances of survival, people with heart disease or those who are at high risk for developing it need early detection and treatment.
4. To pinpoint the essential characteristics of cardiac disease so that patients can get a quick diagnosis and help.

Project Overview

Information on patients is included in the dataset, which has around 300000 records and 18 attributes. The following characteristics are listed among the attributes: HeartDisease, BMI, Smoking, AlcoholDrinking, Stroke, PhysicalHealth, MentalHealth, DiffWalking, Sex, AgeCategory, Race, Diabetes, PhysicalActivity, GenHealth, SleepTime, Asthma, KidneyDisease, SkinCancer. The variable "HeartDisease" is the target variable which is binary ("Yes": The patient has a heart disease; "No": The patient does not have a heart disease), however the classes are not balanced. The data collection is created as a data frame using the Pandas package in Python. Pre-processing and experiments are then conducted using this dataset.

| | HeartDisease | BMI | Smoking | AlcoholDrinking | Stroke | PhysicalHealth | MentalHealth | DiffWalking | Sex | AgeCategory | Race |
|---|--------------|-------|---------|-----------------|--------|----------------|--------------|-------------|--------|-------------|-------|
| 0 | No | 16.60 | Yes | No | No | 3.0 | 30.0 | No | Female | 55-59 | White |
| 1 | No | 20.34 | No | No | Yes | 0.0 | 0.0 | No | Female | 80 or older | White |
| 2 | No | 26.58 | Yes | No | No | 20.0 | 30.0 | No | Male | 65-69 | White |
| 3 | No | 24.21 | No | No | No | 0.0 | 0.0 | No | Female | 75-79 | White |
| 4 | No | 23.71 | No | No | No | 28.0 | 0.0 | Yes | Female | 40-44 | White |

The features

- HeartDisease → Binary (Yes or No)
- BMI (Body mass Index) → (values)
- Smoking → Have you smoked at least 100 cigarettes in your entire life? (The answer is binary (Yes or No))
- AlcoholDrinking → Heavy drinkers: adult men having more than 14 drinks per week and adult women having more than 7 drinks per week (The answer is binary (Yes or No))
- Stroke → (Ever told) (you had) a stroke? (The answer is binary (Yes or No))

- PhysicalHealth → Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good? (0-30 days).
- MentalHealth → Thinking about your mental health, for how many days during the past 30 days was your mental health not good? (0-30 days).
- DiffWalking → Do you have serious difficulty walking or climbing stairs? (The answer is binary (Yes or No))
- Sex → Are you male or female? (The response is binary (Female or Male))
- AgeCategory → Fourteen-level age category.
- Race → Imputed race/ethnicity value.
- Diabetic → (Ever told) (you had) diabetes?
- PhysicalActivity → Adults who reported doing physical activity or exercise during the past 30 days other than their regular job.
- GenHealth → Would you say that in general your health is...
- SleepTime → On average, how many hours of sleep do you get in a 24-hour period?
- Asthma → (Ever told) (you had) asthma?
- KidneyDisease → Not including kidney stones, bladder infection or incontinence, were you ever told you had kidney disease?
- SkinCancer → (Ever told) (you had) skin cancer?

Categorical Features: HeartDisease, Smoking, AlcoholDrinking, Stroke, DiffWalking, Sex, Race, Diabetic, PhysicalActivity, GenHealth, Asthma, KidneyDisease, SkinCancer

Continuous Features: BMI, PhysicalHealth, MentalHealth, AgeCategory, SleepTime

Workspace

Importing file using pandas

```
[ ] heart = pd.read_csv('heart_2020_cleaned.csv')
```

```
[ ] original = heart
```

```
[ ] heart.head()
```

| | HeartDisease | BMI | Smoking | AlcoholDrinking | Stroke | PhysicalHealth | MentalHealth | DiffWalking | Sex | AgeCategory | Race | Diabetic | PhysicalActivity |
|---|--------------|-------|---------|-----------------|--------|----------------|--------------|-------------|--------|-------------|-------|----------|------------------|
| 0 | No | 16.60 | Yes | No | No | 3.0 | 30.0 | No | Female | 55-59 | White | Yes | Yes |
| 1 | No | 20.34 | No | No | Yes | 0.0 | 0.0 | No | Female | 80 or older | White | No | Yes |
| 2 | No | 26.58 | Yes | No | No | 20.0 | 30.0 | No | Male | 65-69 | White | Yes | Yes |
| 3 | No | 24.21 | No | No | No | 0.0 | 0.0 | No | Female | 75-79 | White | No | No |
| 4 | No | 23.71 | No | No | No | 28.0 | 0.0 | Yes | Female | 40-44 | White | No | Yes |

Dataset information

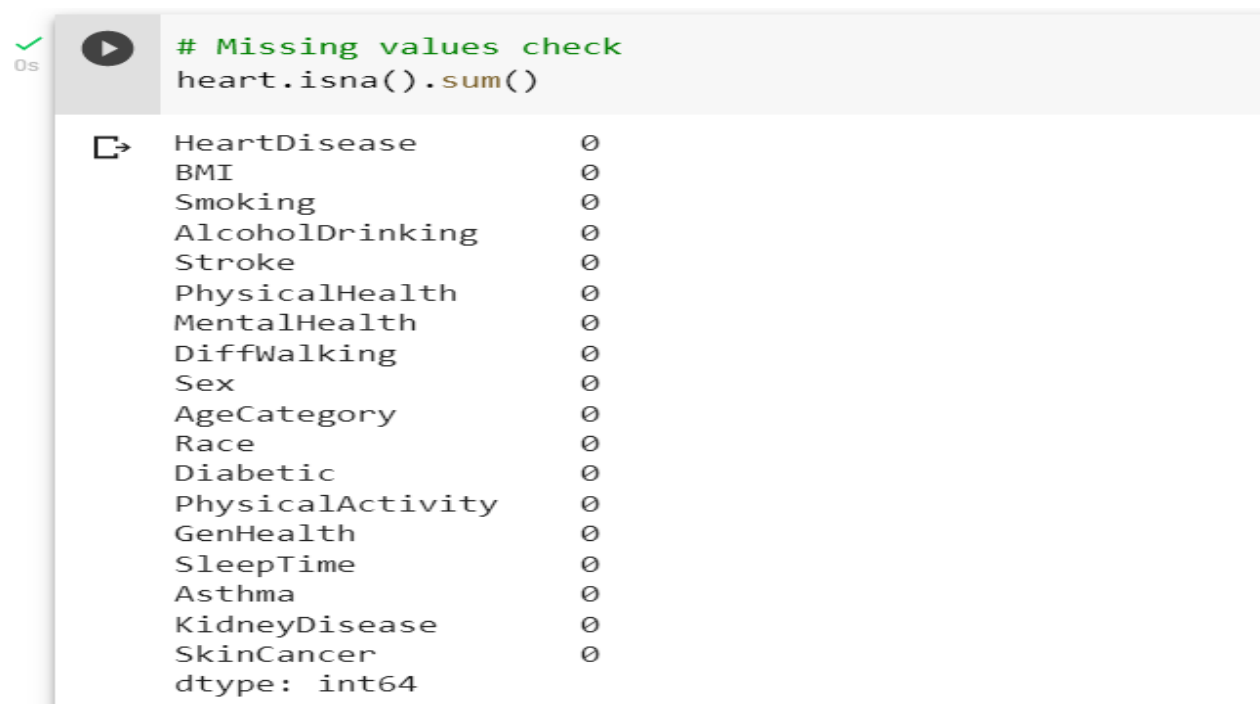
```
[ ] heart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 319795 entries, 0 to 319794
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   HeartDisease          319795 non-null object
1   BMI                   319795 non-null float64
2   Smoking               319795 non-null object
3   AlcoholDrinking       319795 non-null object
4   Stroke                319795 non-null object
5   PhysicalHealth         319795 non-null float64
6   MentalHealth          319795 non-null float64
7   DiffWalking           319795 non-null object
8   Sex                   319795 non-null object
9   AgeCategory           319795 non-null object
10  Race                  319795 non-null object
11  Diabetic              319795 non-null object
12  PhysicalActivity       319795 non-null object
13  GenHealth             319795 non-null object
14  SleepTime             319795 non-null float64
15  Asthma                319795 non-null object
16  KidneyDisease          319795 non-null object
17  SkinCancer            319795 non-null object
dtypes: float64(4), object(14)
memory usage: 43.9+ MB
```

Exploratory Data Analysis

Checking missing values

To properly handle the remaining data, it is crucial to determine whether any values are missing from the data. Even if some values are randomly absent, the data sample nevertheless probably represents the population. But the analysis might be skewed if the values are absent systematically.



```
# Missing values check
heart.isna().sum()
```

| | |
|------------------|---|
| HeartDisease | 0 |
| BMI | 0 |
| Smoking | 0 |
| AlcoholDrinking | 0 |
| Stroke | 0 |
| PhysicalHealth | 0 |
| MentalHealth | 0 |
| DiffWalking | 0 |
| Sex | 0 |
| AgeCategory | 0 |
| Race | 0 |
| Diabetic | 0 |
| PhysicalActivity | 0 |
| GenHealth | 0 |
| SleepTime | 0 |
| Asthma | 0 |
| KidneyDisease | 0 |
| SkinCancer | 0 |
| dtype: int64 | |

Checking unique values

```

▶ heart.nunique()

HeartDisease      2
BMI              3604
Smoking           2
AlcoholDrinking   2
Stroke            2
PhysicalHealth    31
MentalHealth      31
DiffWalking       2
Sex               2
AgeCategory       13
Race              6
Diabetic          4
PhysicalActivity   2
GenHealth         5
SleepTime         24
Asthma            2
KidneyDisease     2
SkinCancer        2
dtype: int64

```


Correlation Analysis

```

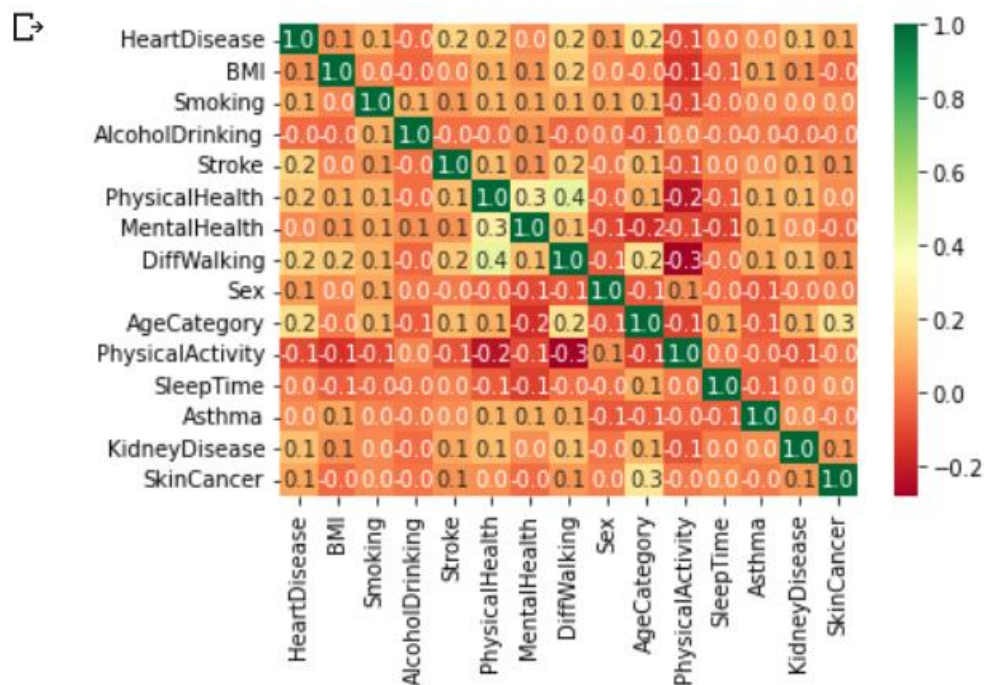
#Correlation

fig = px.imshow(heart[['HeartDisease', 'Smoking', 'AlcoholDrinking', 'Stroke', 'DiffWalking', 'Sex',
                        'PhysicalActivity', 'Asthma', 'KidneyDisease', 'SkinCancer', 'BMI', 'PhysicalHealth',
                        'MentalHealth', 'AgeCategory', 'SleepTime']].corr(), color_continuous_scale="Blues")
fig.update_layout(height=800)
fig.show()

```



Correlation Heatmap



No strong correlation has been found among the variables.

Data Distribution



Column stats for continuous data







```
heart.describe()[1:][['BMI', 'PhysicalHealth', 'MentalHealth', 'AgeCategory', 'SleepTime']].T.style.background_gradient(cmap='Blues')
```


| | mean | std | min | 25% | 50% | 75% | max |
|-----------------------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| BMI | 28.325399 | 6.356100 | 12.020000 | 24.030000 | 27.340000 | 31.420000 | 94.850000 |
| PhysicalHealth | 3.371710 | 7.950850 | 0.000000 | 0.000000 | 0.000000 | 2.000000 | 30.000000 |
| MentalHealth | 3.898366 | 7.955235 | 0.000000 | 0.000000 | 0.000000 | 3.000000 | 30.000000 |
| AgeCategory | 54.355759 | 17.720429 | 21.000000 | 42.000000 | 57.000000 | 67.000000 | 80.000000 |
| SleepTime | 7.097075 | 1.436007 | 1.000000 | 6.000000 | 7.000000 | 8.000000 | 24.000000 |

Column stats for categorical data

```
fig = make_subplots(rows=7, cols=2, subplot_titles=("HeartDisease", "Smoking", "AlcoholDrinking", "Stroke",
"DiffWalking", "Sex", "Race", "Diabetic", "PhysicalActivity", "GenHealth", "Asthma",
"KidneyDisease", "SkinCancer"),
specs=[[{"type": "domain"}, {"type": "domain"}], [{"type": "domain"}, {"type": "domain"}],
[{"type": "domain"}, {"type": "domain"}], [{"type": "domain"}, {"type": "domain"}],
[{"type": "domain"}, {"type": "domain"}], [{"type": "domain"}, {"type": "domain"}],
[{"type": "domain"}, {"type": "domain"}], [{"type": "domain"}, {"type": "domain"}],])
```

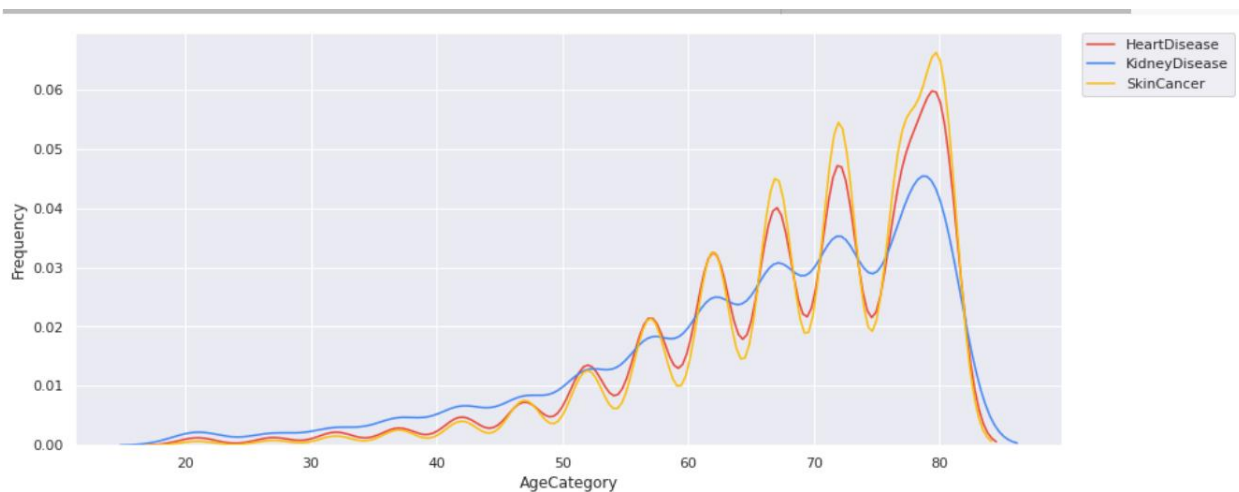





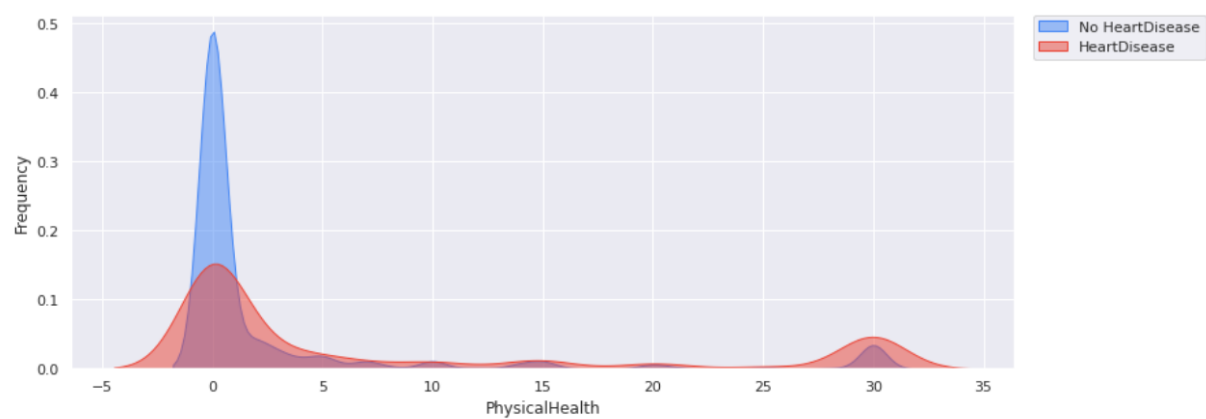
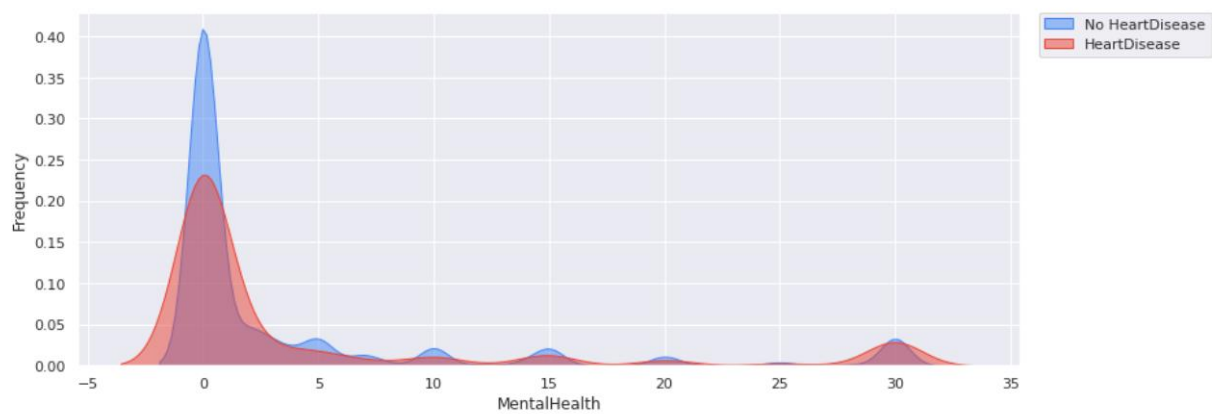
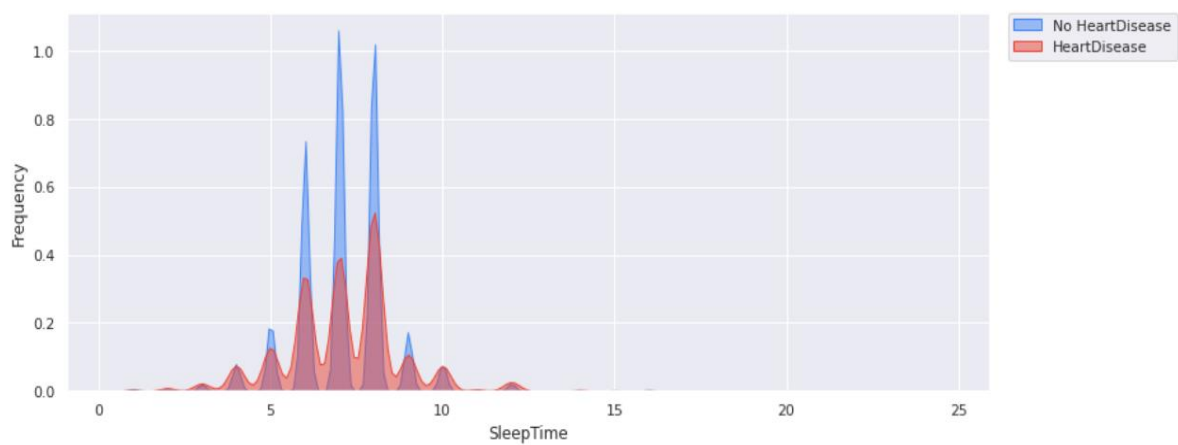



The target variable (HeartDisease) has an uneven distribution, as seen in the plots above, so the data needs to be balanced before developing the model.

Data Visualization



Insights - People having heart disease, skin cancer and kidney disease are mostly old people






```
female_with_heart_disease = len(heart[(heart['HeartDisease']=="Yes") & (heart['Sex']=="Female")])
num_female = len(heart[heart['Sex']=="Female"])
male_with_heart_disease = len(heart[(heart['HeartDisease']=="Yes") & (heart['Sex']=="Male")])
num_male = len(heart[heart['Sex']=="Male"])
print('Probability of Male to have Heart disease:', male_with_heart_disease/num_male)
print('Probability of Female to have Heart disease:', female_with_heart_disease/num_female)
```

```
Probability of Male to have Heart disease: 0.10618461740904007
Probability of Female to have Heart disease: 0.06694675367241738
```

Insights -

1. Mostly heart disease patients are Males than Females
2. More Females were tested than Males
3. Males are approximately 1.6 times more likely to have heart disease than females



```

smoke_and_heart_disease = len(heart[(heart['HeartDisease']=="Yes") & (heart['Smoking']=="Yes")])
num_smoke = len(heart[heart['Smoking']=="Yes"])
no_smoke_and_heart_disease = len(heart[(heart['HeartDisease']=="Yes") & (heart['Smoking']=="No")])
num_no_smoke = len(heart[heart['Smoking']=="No"])
print('Probability of Heart disease if you smoke:', smoke_and_heart_disease/num_smoke)
print("Probability of Heart disease if you don't smoke:", no_smoke_and_heart_disease/num_no_smoke)

```

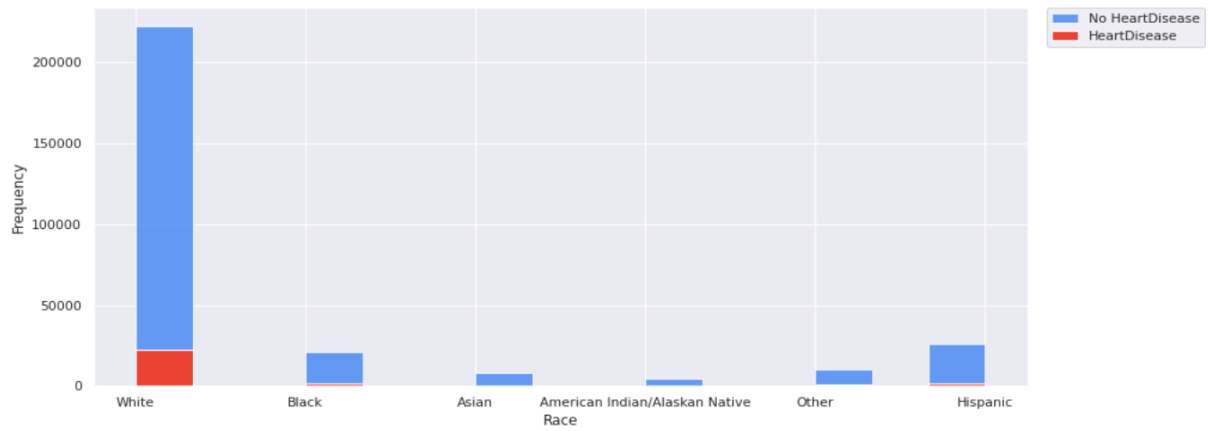
```

Probability of Heart disease if you smoke: 0.12157715983867544
Probability of Heart disease if you don't smoke: 0.0603341370078824

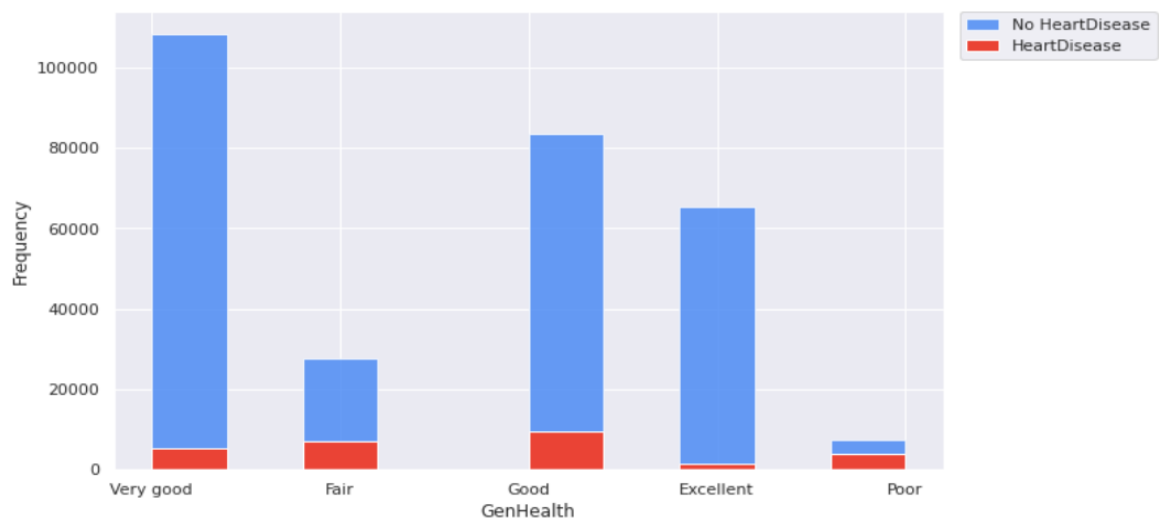
```

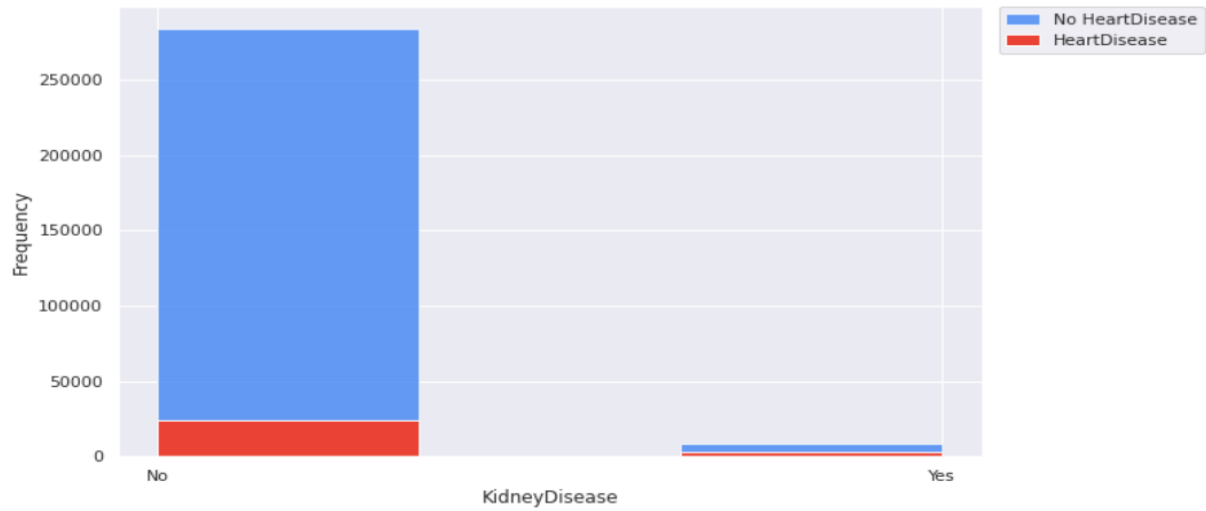
Insights -

1. Most heart disease patients smoke.
2. People who smoke are approximately twice as likely to have heart disease than people who don't smoke.

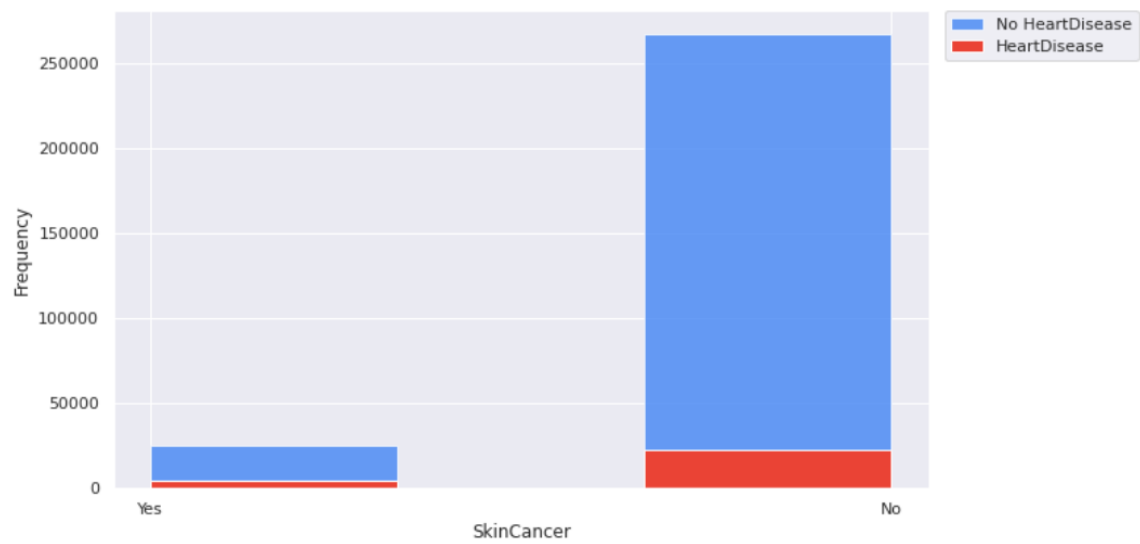



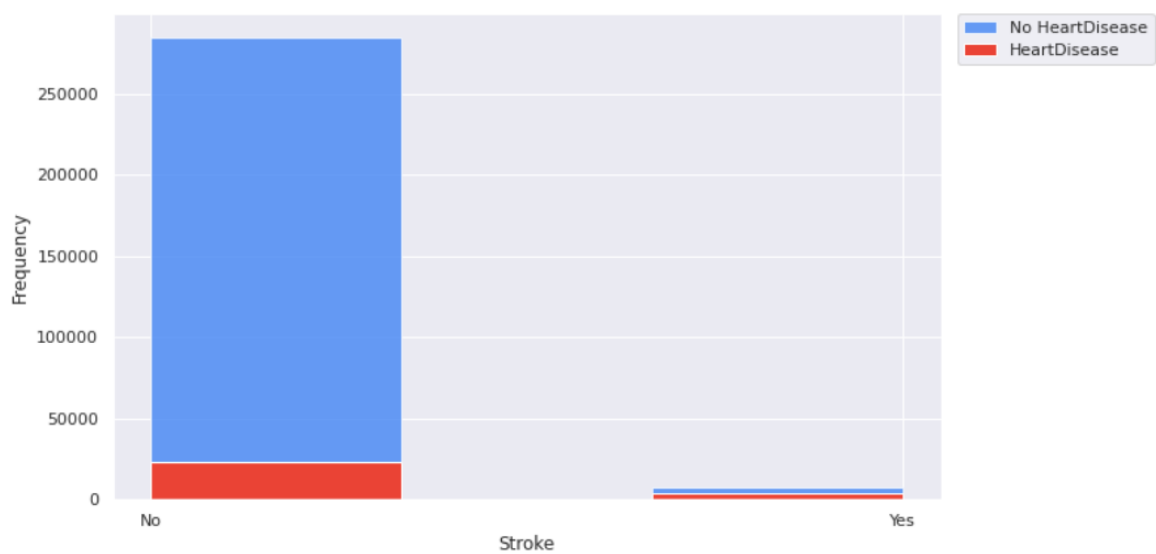
Insights - Most heart disease patients are white people.


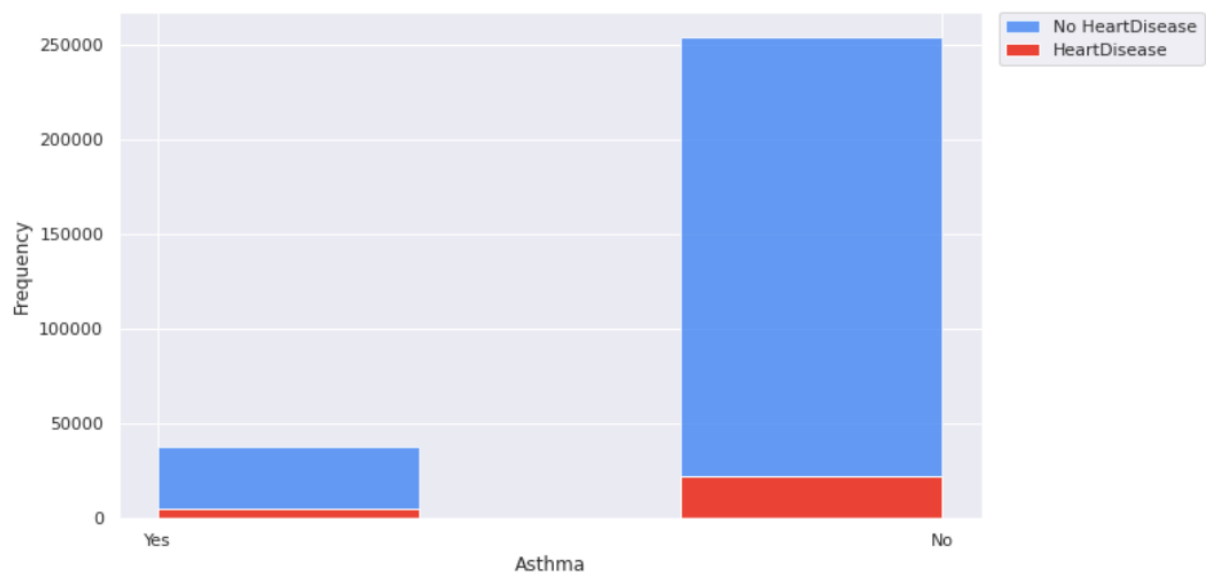


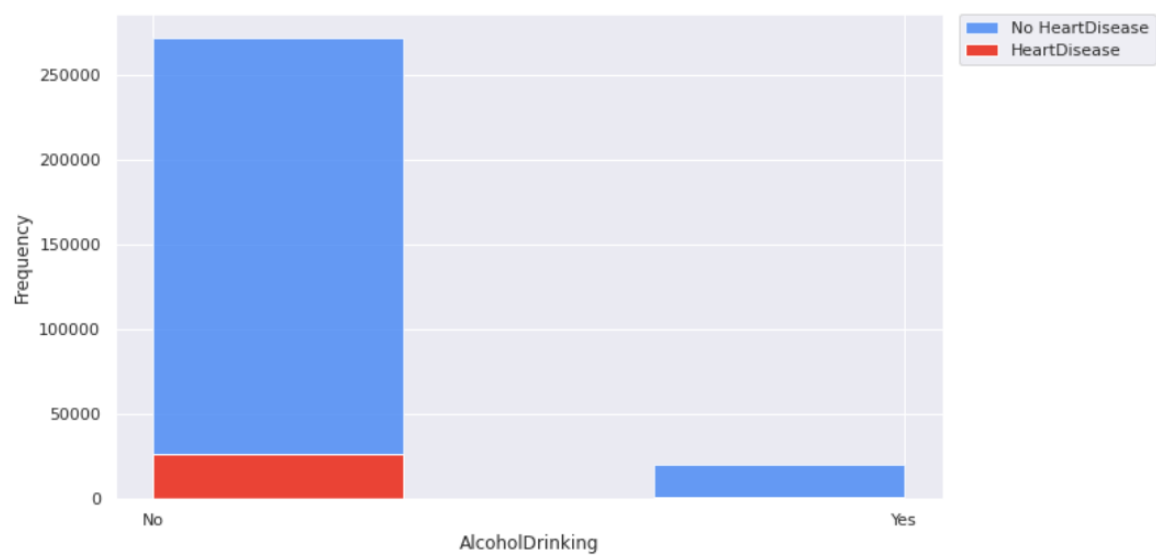
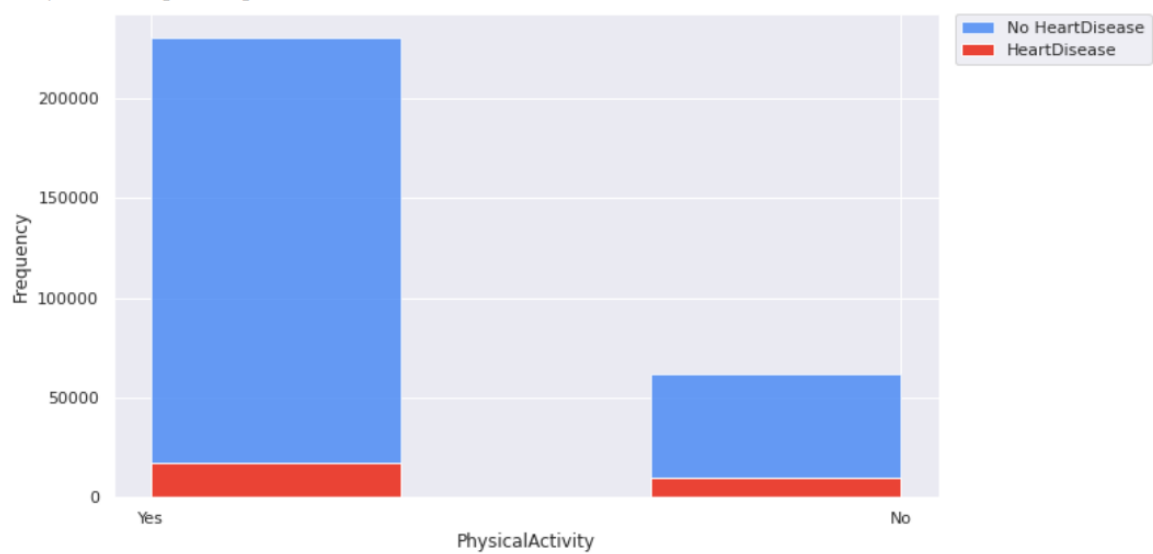


Insights - Most of the people tested have no kidney disease.









Data Cleaning and Preprocessing

```
print('Continuous Columns')
heart.select_dtypes(include=['float']).describe().T[['min', 'mean', 'max']].style.background_gradient(cmap='Blues')
```

Continuous Columns

| | min | mean | max |
|-----------------------|-----------|-----------|-----------|
| BMI | 0.126726 | 0.298634 | 1.000000 |
| PhysicalHealth | 0.000000 | 0.112390 | 1.000000 |
| MentalHealth | 0.000000 | 0.129946 | 1.000000 |
| AgeCategory | 21.000000 | 54.355759 | 80.000000 |
| SleepTime | 1.000000 | 7.097075 | 24.000000 |

The range of continuous features is different.

So, dividing them by the highest value of each column to scale them to be between 0 and 1.

```
for col in ['BMI', 'PhysicalHealth', 'MentalHealth', 'AgeCategory', 'SleepTime']:
    heart[col] = heart[col]/heart[col].max()
```

```
print('\nCategorical Columns\n')
heart.select_dtypes(include=['O']).nunique()
```

Categorical Columns

```
HeartDisease      2
Smoking           2
AlcoholDrinking   2
Stroke            2
DiffWalking       2
Sex               2
Race              6
Diabetic          4
PhysicalActivity   2
GenHealth         5
Asthma            2
KidneyDisease     2
SkinCancer        2
dtype: int64
```

There are categories columns with only two values and others with multiple values. Here, the columns with two distinct values will be converted to binary (either 1 or 0), while the other categorical columns with more than two distinct values will be one-hot encoded.

```
# Integer encode columns with 2 unique values
for col in ['HeartDisease', 'Smoking', 'AlcoholDrinking', 'Stroke', 'DiffWalking', 'Sex', 'PhysicalActivity', 'Asthma', 'KidneyDi
    if heart[col].dtype == 'O':
        le = LabelEncoder()
        heart[col] = le.fit_transform(heart[col])
# One-hot encode columns with more than 2 unique values
df = pd.get_dummies(heart, columns=['Race', 'Diabetic', 'GenHealth'], prefix = ['Race', 'Diabetic', 'GenHealth'])
```

Splitting the dataset

```
#Split dataset

X = heart.drop(columns = ['HeartDisease'])
y = heart['HeartDisease']
X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size = 0.4, random_state = 1)
```

As mentioned earlier, the dataset is not balanced, there are a greater number of samples for people who have no heart disease than people with heart disease. So, this can be fixed by oversampling the data using SMOTE technique.

```
sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_resample(X, y)

df_smote_over = pd.concat([pd.DataFrame(X_res), pd.DataFrame(y_res)], axis = 1)
```

```
y_res.value_counts()
```

```
0    292422
1    292422
Name: HeartDisease, dtype: int64
```

```
X = df_smote_over.drop(columns = ['HeartDisease'])
y = df_smote_over['HeartDisease']
X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size = 0.4, random_state = 1)
```


One of the most popular oversampling techniques to address the imbalance issue is SMOTE (synthetic minority oversampling technique). By increasing the minority class samples at random and duplicating them, it seeks to balance the distribution of classes.

So, now the data is balanced, and modelling can be performed.

Model building

Random Forest (variable/feature importance)

A classification system made up of several decision trees is called the random forest. It attempts to produce an uncorrelated forest of trees whose forecast by committee is more accurate than that of any individual tree by using bagging and feature randomness when generating each individual tree.

Knowing the feature importance suggested by machine learning models can be useful in a variety of ways, such as:

- Understanding the model's logic will enable you to not only confirm that it is accurate but also to work on improving the model by concentrating just on the crucial variables.
- It can be used for variable selection, allowing you to eliminate x variables that have little bearing on performance and require significantly less training time.
- In certain commercial situations, it makes sense to give up some accuracy in favour of interpretability. For instance, if a bank denies a loan application, it must provide the customer with the justification for the decision.

Random forest

```
[ ] random_forest = RandomForestClassifier(random_state=1, n_estimators=500)
    random_forest.fit(X_train, y_train)
```

```
RandomForestClassifier(n_estimators=500, random_state=1)
```

```
[ ] importance = random_forest.feature_importances_
    std = np.std([tree.feature_importances_ for tree in random_forest.estimators_], axis=0)
```

Calculating the importance of each variable for randomforest

```
[ ] random_forest_df = pd.DataFrame({"feature" : X_train.columns,
                                     "importance" : importance,
                                     "Std" : std})
    print(random_forest_df.sort_values("importance", ascending = False))
```

| | feature | importance | Std |
|----|----------------------------------|------------|----------|
| 13 | AgeCategory | 0.286817 | 0.052090 |
| 0 | BMI | 0.182437 | 0.005123 |
| 9 | SleepTime | 0.138979 | 0.020963 |
| 4 | PhysicalHealth | 0.077434 | 0.024123 |
| 5 | MentalHealth | 0.050789 | 0.006052 |
| 6 | DiffWalking | 0.031618 | 0.023939 |
| 20 | Diabetic_Yes | 0.027975 | 0.017517 |
| 7 | Sex | 0.023321 | 0.002525 |
| 22 | GenHealth_Fair | 0.020681 | 0.011942 |
| 3 | Stroke | 0.018405 | 0.009291 |
| 1 | Smoking | 0.018004 | 0.005296 |
| 23 | GenHealth_Good | 0.016156 | 0.004574 |
| 8 | PhysicalActivity | 0.015296 | 0.003228 |
| 24 | GenHealth_Poor | 0.013897 | 0.011526 |
| 25 | GenHealth_Very good | 0.012543 | 0.006467 |
| 12 | SkinCancer | 0.011360 | 0.002479 |
| 10 | Asthma | 0.011273 | 0.000763 |
| 18 | Race_White | 0.008943 | 0.001239 |
| 11 | KidneyDisease | 0.008178 | 0.003845 |
| 2 | AlcoholDrinking | 0.006380 | 0.000631 |
| 15 | Race_Black | 0.005058 | 0.000503 |
| 16 | Race_Hispanic | 0.004745 | 0.000647 |
| 19 | Diabetic_No, borderline diabetes | 0.003832 | 0.000323 |
| 17 | Race_Other | 0.003276 | 0.000321 |
| 14 | Race_Asian | 0.001843 | 0.000233 |
| 21 | Diabetic_Yes (during pregnancy) | 0.000758 | 0.000139 |

```
classificationSummary(y_train, random_forest.predict(X_train))
```

Confusion Matrix (Accuracy 0.9985)

| | Prediction | |
|--------|------------|--------|
| Actual | 0 | 1 |
| 0 | 175332 | 259 |
| 1 | 257 | 175058 |

```
classificationSummary(y_valid, random_forest.predict(X_valid))
```

Confusion Matrix (Accuracy 0.9200)

| | Prediction | |
|--------|------------|--------|
| Actual | 0 | 1 |
| 0 | 106900 | 9931 |
| 1 | 8794 | 108313 |

Random Forest Classifier

Random Forest

```
[ ] model=RandomForestClassifier()
```

```
[ ] model.fit(X_train, y_train)
```


```
RandomForestClassifier()
```

```
[ ] rfc_result = evaluate_model(model, X_valid, y_valid)
```

```
[ ] print_model_result(rfc_result)
```

```
Accuracy: 0.9194743906505143
Precision: 0.916214178617716
Recall: 0.9235997848121803
F1 Score: 0.9198921576131792
Area Under Curve: 0.9696814367295973
Confusion Matrix:
[[106940  9891]
 [ 8947 108160]]
```

```
[ ] draw_cm("RandomForestClassifier", rfc_result["cm"])
```



Decision Tree

Making decisions with decision trees is effective because they clearly state the problem so that all alternative solutions may be tested. Decision tree allows us to carefully consider any possible consequences of a decision. It creates a framework for calculating outcome values and success probabilities.

```
#Decision Tree

model=tree.DecisionTreeClassifier()

model.fit(X_train, y_train)


DecisionTreeClassifier()

dt_result = evaluate_model(model, X_valid, y_valid)

print_model_result(dt_result)

Accuracy: 0.8910950764732536
Precision: 0.8887700897781851
Recall: 0.8943786451706559
F1 Score: 0.8915655471519836
Area Under Curve: 0.8917678098860559
Confusion Matrix:
[[103723  13108]
 [ 12369 104738]]

draw_cm("DecisionTreeClassifier",dt_result["cm"])
```



Grid search

```
[ ] param_grid = {
    'max_depth' : [2,3,5,10],
    'min_samples_split' : [0.07,0.05,0.01,0.005],
    'min_impurity_decrease' : [0.05, 0.02, 0.01, 0.001]
}
```

```
[ ] gridsearch = GridSearchCV(DecisionTreeClassifier(random_state=1), param_grid, cv=5, n_jobs= -1)
gridsearch.fit(X_train,y_train)
```

```
GridSearchCV(cv=5, estimator=DecisionTreeClassifier(random_state=1), n_jobs=-1,
             param_grid={'max_depth': [2, 3, 5, 10],
                         'min_impurity_decrease': [0.05, 0.02, 0.01, 0.001],
                         'min_samples_split': [0.07, 0.05, 0.01, 0.005]})
```

```
[ ] gridsearch.best_params_
```

```
{'max_depth': 10, 'min_impurity_decrease': 0.001, 'min_samples_split': 0.01}
```

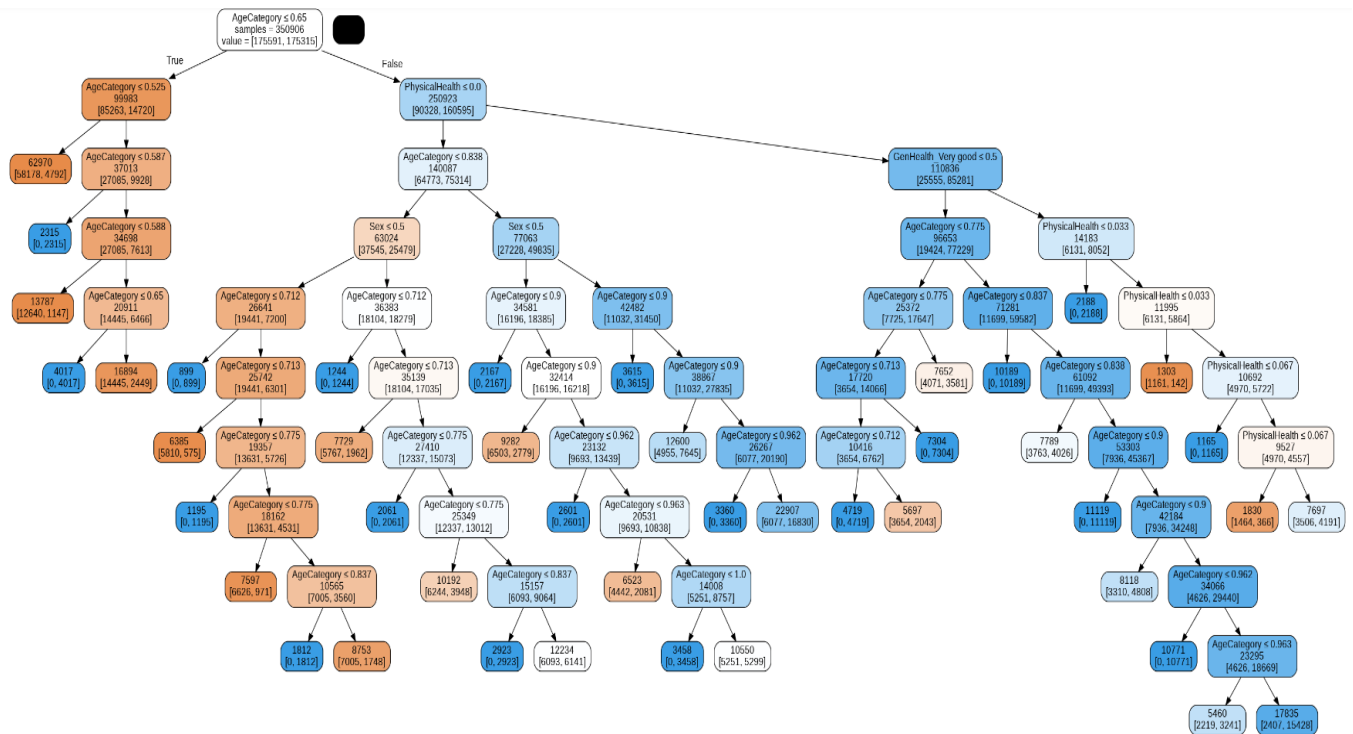
```
[ ] gridsearch.best_score_
```

```
0.8097467709940634
```

```
gridsearch.best_estimator_
```

```
DecisionTreeClassifier(max_depth=10, min_impurity_decrease=0.001,
                       min_samples_split=0.01, random_state=1)
```

```
gridsearchtree = gridsearch.best_estimator_
plotDecisionTree(gridsearchtree, feature_names = X_train.columns)
```



```
regressionSummary(y_train, gridsearchtree.predict(X_train))
```

Regression statistics

Mean Error (ME) : -0.0256
 Root Mean Squared Error (RMSE) : 0.4342
 Mean Absolute Error (MAE) : 0.1886

```
regressionSummary(y_valid, gridsearchtree.predict(X_valid))
```

Regression statistics

Mean Error (ME) : -0.0252
 Root Mean Squared Error (RMSE) : 0.4325
 Mean Absolute Error (MAE) : 0.1870

```
classificationSummary(y_train, gridsearchtree.predict(X_train))
```

Confusion Matrix (Accuracy 0.8114)

| | Prediction | |
|--------|------------|--------|
| Actual | 0 | 1 |
| 0 | 138010 | 37581 |
| 1 | 28584 | 146731 |

```
classificationSummary(y_valid, gridsearchtree.predict(X_valid))
```

Confusion Matrix (Accuracy 0.8130)

| | Prediction | |
|--------|------------|-------|
| Actual | 0 | 1 |
| 0 | 92002 | 24829 |
| 1 | 18924 | 98183 |

Logistic Regression

Based on previous observations of a data set, the statistical analysis technique of logistic regression can be used to forecast a binary outcome, such as yes or no. A logistic regression model uses an analysis of the correlation between one or more pre-existing independent variables to predict a dependent data variable.

The technique of logistic regression has grown in significance in the field of machine learning. It enables machine learning algorithms to categorise incoming input based on previous data. The algorithms get more accurate at predicting classes within data sets when new pertinent data is added.

```
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
```

```
logit_reg = LogisticRegression(solver='liblinear', C=1e42, random_state=1)
logit_reg.fit(X_train, y_train)
```

```
LogisticRegression(C=1e+42, random_state=1, solver='liblinear')
```



```
print(pd.DataFrame({'coef': logit_reg.coef_[0]}, index=X.columns))
```

| | coef |
|----------------------------------|-----------|
| BMI | 0.784984 |
| Smoking | 0.370586 |
| AlcoholDrinking | -0.494361 |
| Stroke | 1.039995 |
| PhysicalHealth | 0.087509 |
| MentalHealth | 0.130238 |
| DiffWalking | 0.209328 |
| Sex | 0.722479 |
| PhysicalActivity | -0.031046 |
| SleepTime | -0.721561 |
| Asthma | 0.275082 |
| KidneyDisease | 0.414194 |
| SkinCancer | 0.045557 |
| AgeCategory | 5.267971 |
| Race_Asian | -0.860504 |
| Race_Black | -0.514394 |
| Race_Hispanic | -0.378614 |
| Race_Other | -0.285395 |
| Race_White | -0.387344 |
| Diabetic_No, borderline diabetes | -0.335998 |
| Diabetic_Yes | 0.443877 |
| Diabetic_Yes (during pregnancy) | -0.976448 |
| GenHealth_Fair | 1.559602 |
| GenHealth_Good | 1.033959 |
| GenHealth_Poor | 1.994703 |
| GenHealth_Very good | 0.428447 |

```
logit_reg_prob = logit_reg.predict_proba(X_valid)
```

```
logit_reg_pred = logit_reg.predict(X_valid)
```

```
logit_result = pd.DataFrame({'actual' : y_valid,
                             'p_0' : [p[0] for p in logit_reg_prob],
                             'p_1' : [p[1] for p in logit_reg_prob],
                             'predicted': logit_reg_pred})
logit_result
```

| | actual | p_0 | p_1 | predicted |
|--------|--------|----------|----------|-----------|
| 348389 | 1 | 0.674225 | 0.325775 | 0 |
| 296815 | 0 | 0.464427 | 0.535573 | 1 |
| 498134 | 1 | 0.549959 | 0.450041 | 0 |
| 538898 | 1 | 0.467675 | 0.532325 | 1 |
| 575047 | 1 | 0.085597 | 0.914403 | 1 |
| ... | ... | ... | ... | ... |
| 477789 | 1 | 0.166636 | 0.833364 | 1 |
| 68364 | 0 | 0.342609 | 0.657391 | 1 |
| 248603 | 0 | 0.217977 | 0.782023 | 1 |
| 392711 | 1 | 0.676344 | 0.323656 | 0 |
| 180660 | 0 | 0.988944 | 0.011056 | 0 |

233938 rows × 4 columns

```
classificationSummary(y_train, logit_reg.predict(X_train))
```

Confusion Matrix (Accuracy 0.7650)

| | | |
|--------|------------|--------|
| | Prediction | |
| Actual | 0 | 1 |
| 0 | 129632 | 45959 |
| 1 | 36504 | 138811 |

```
classificationSummary(y_valid, logit_reg.predict(X_valid))
```

Confusion Matrix (Accuracy 0.7657)

| | | |
|--------|------------|-------|
| | Prediction | |
| Actual | 0 | 1 |
| 0 | 86364 | 30467 |
| 1 | 24339 | 92768 |

Logistic Regression

```
[ ] model=sklin.LogisticRegression()

[ ] model.fit(X_train, y_train)

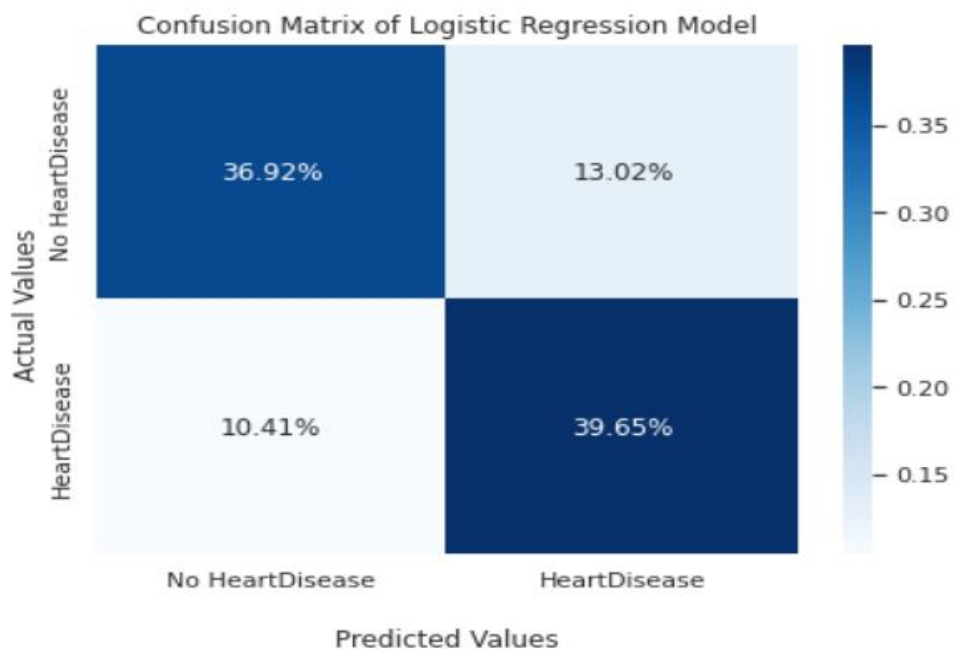
LogisticRegression()

[ ] logistic_regression_result = evaluate_model(model, X_valid, y_valid)

[ ] print_model_result(logistic_regression_result)

Accuracy: 0.7657199770879464
Precision: 0.7527793556763775
Recall: 0.792138813222096
F1 Score: 0.7719577093830747
Area Under Curve: 0.8414418540801041
Confusion Matrix:
[[86366 30465]
 [24342 92765]]

[ ] draw_cm("Logistic Regression",logistic_regression_result["cm"])
```



XG Boost

XG Boost

```

▶ from xgboost import XGBClassifier
  model=XGBClassifier()

[ ] model.fit(X_train, y_train)

    XGBClassifier()


[ ] xgb_result = evaluate_model(model, X_valid, y_valid)

[ ] print_model_result(xgb_result)

Accuracy: 0.8710213817336218
Precision: 0.8803352991617521
Recall: 0.8591288308982383
F1 Score: 0.869602796973115
Area Under Curve: 0.9520035868653651
Confusion Matrix:
[[103155  13676]
 [ 16497 100610]]


[ ] draw_cm("XGBClassifier",dt_result["cm"])


```



Model Comparison (Best Model)

Among all the models, Random Forest is the best model with highest accuracy. Also, models are being evaluated on various metrics like Accuracy, Precision, Recall and F1 score. Along with accuracy, the best model is considered based on the ROC curve value.





Interpretation and Conclusion

- Adults can work on improving their overall health.
- The BMI affects heart disease. Heart disease tends to be more likely to affect people with a BMI of 20 or above.
- Reduce chances of having heart disease by doing physical activity and eating healthy food.
- Smoking causes heart disease, among other things. Contrary to popular belief, the data indicates that alcohol consumption does not increase the risk of heart disease.

- Heart disease is influenced by the number of physical health issues faced throughout the previous 30 days. On the other hand, the research indicates that the incidence of mental health issues during the same time only contributes marginally to heart disease.
- Stroke contributes to heart disease.
- Heart disease is more likely to affect those who have trouble walking.
- In terms of sexes, men typically have a larger risk of developing heart disease than women.
- The amount of physical exercise that was not work-related throughout the previous 30 days reduces the risk of developing heart disease. The general health conditions come to the same conclusion.
- A day's worth of sleep does not appear to have a significant impact on the development of heart disease. However, those who are heart disease-free often get roughly 6 to 8 hours of better sleep per day.
- Heart disease is influenced by asthma, renal illness, skin cancer, and diabetes. Asthma, however, only makes a minor impact.
- Heart disease is more common in persons over 60 than in those under 60.
- It does not appear that race affects the risk of heart disease. Any race could be affected by heart disease.

Recommendations

- People should routinely have their heart checkup. Even if a person is not currently dealing with any health issues, the person could still be in danger. Annual checkups should include a heart disease screening.
- For balancing blood sugar and cholesterol levels, doctors should emphasise the significance of a healthy diet and exercise. People should refrain from binge drinking and smoking since these vices raise the risk of developing heart disease.
- The likelihood of developing heart disease is significantly increased in anyone who already experienced a heart attack. Heart attack survivors should have routine screenings to look for any abnormalities that could be signs of heart disease.